

# Image Search with Python

Arif Qodari



Photo by Toa Heftiba

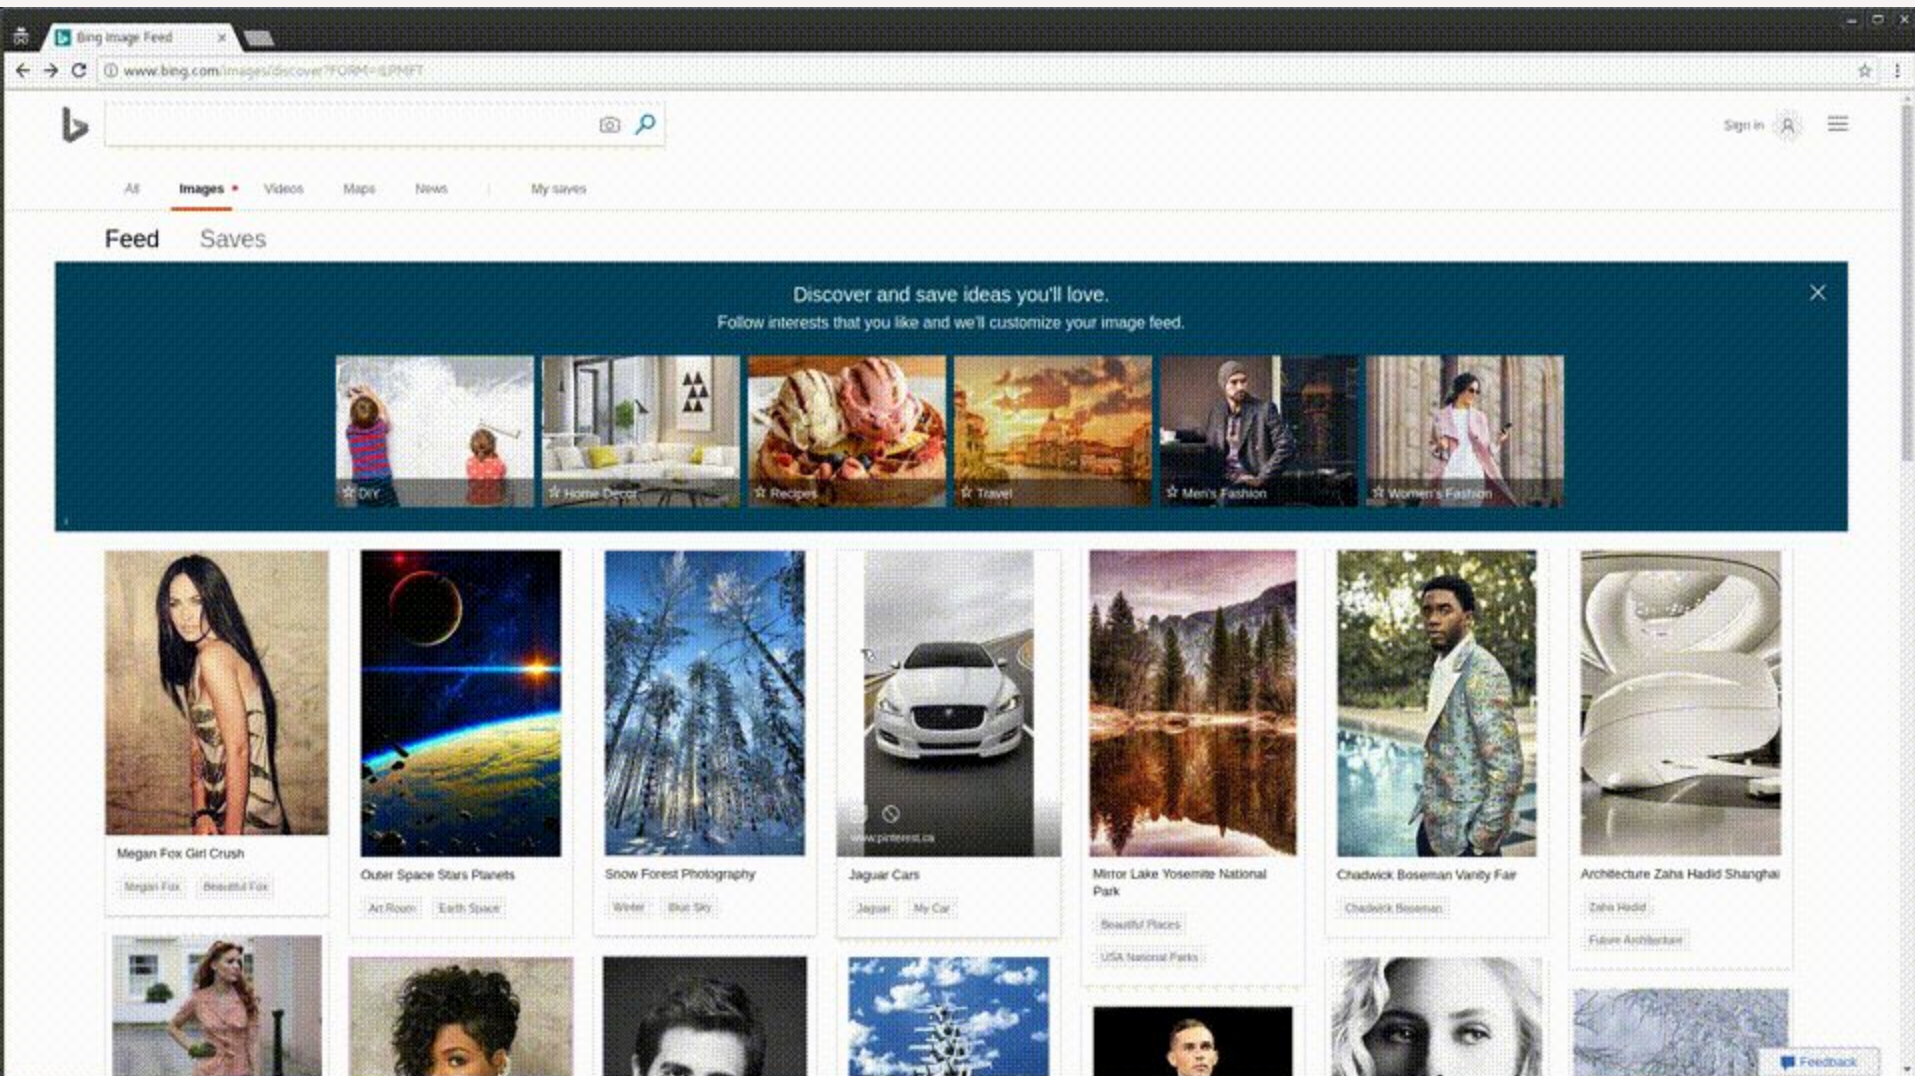
# About me

- Currently Data Scientist at [salestock.id](https://salestock.id)
- Previously Ruby on Rails and WordPress dev
- Python and (sometimes) C++
- @arifgodari on Twitter and Github

**What is image search engine?**



# “I have an image, give me similar images.”



**How to make images searchable?**

Let's break down into 2 questions:

- 1. What kind of features to be used?**
- 2. How to search similar features?**

# Visual Features

1. **Color composition**, e.g. histogram of color
2. **Shape of objects**, e.g. contour
3. **Texture**, e.g. local binary pattern
4. **Combinations ??**
5. **Deep Learned Features ??**

## Basic features extraction pipeline




Features  
extraction

[0.13, 0.91, 0.62, ... ,0.71, 0.42]



# 1. Histogram of Color (HoC)



# How to Construct HoC

- Typically use HSV color instead of RGB
- Define number of bins (or do some experiments!)
- Count number of pixels for each bin in each channel

```
import imageio
import numpy as np
from skimage.color import rgb2hsv

image = imageio.imread("nasipadang.jpg")
hsv_image = rgb2hsv(image)
(hist, bins) = np.histogram(hsv_image[:, :, 0], bins=32)
```

## 5. FC6 Layer Output as Features




Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Let's break down into 2 questions:

- 1. What kind of features to be used? ✓**
- 2. How to search similar features?**

# **(Dis)similarity Metric**

- **Dissimilarity**, e.g. Euclidean distance
- **Similarity**, e.g. Cosine similarity

[0.13, 0.91, 0.62, ... ,0.71, 0.42]



(dis)similarity  
metric

0.87

[0.13, 0.85, 0.62, ... ,0.71, 0.42]

$$\begin{aligned}d(\mathbf{p}, \mathbf{q}) &= d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \cdots + (q_n - p_n)^2} \\&= \sqrt{\sum_{i=1}^n (q_i - p_i)^2}.\end{aligned}$$

```
import numpy as np

# euclidean distance between vector A and B
# the smaller distance the more similar images
np.linalg.norm(b - a)
```



$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

```
import numpy as np

# cosine similarity between vector A and B
# the higher the value the more similar images
np.dot(a, b) / (np.linalg.norm(a) * np.linalg.norm(b))
```

**... then how to search similar images?**

# (1) Exhaustive Search

Imagine performing a linear scan ( $O(n)$ ) on your MySQL table every time someone hit the search button.


image_id	f1	f2	f3	f4	f5
1	0.340	0.329	0.871	0.261	0.730
2	0.041	0.042	0.570	0.833	0.569
3	0.924	0.966	0.453	0.154	0.846
4	0.534	0.499	0.006	0.555	0.369
5	0.277	0.594	0.753	0.540	0.736
6	0.898	0.554	0.241	0.993	0.353
7	0.019	0.883	0.065	0.795	0.676
8	0.643	0.906	1.000	0.901	0.954
9	0.314	0.399	0.953	0.493	0.851
10	0.517	0.914	0.668	0.101	0.101

```
-- input [0.340 0.329 0.871 0.261 0.730]
SELECT
    image_id,
    SQRT(POWER((0.340 - f1), 2)
        + POWER((0.329 - f2), 2)
        + POWER((0.871 - f3), 2)
        + POWER((0.261 - f4), 2)
        + POWER((0.730 - f5), 2))
    AS distance
FROM db_features
ORDER BY 2
```

**It works but ...**



## (2) Better Approach: Nearest Neighbor



# annoy

Library for approximate nearest neighbor. Written in C++ with Python bindings. Based on forest of binary trees.


<https://github.com/spotify/annoy/>


**Alternatives:** [kd-tree](#), [KGraph](#), [Faiss](#), [NMSLib](#)

# Start with the point set










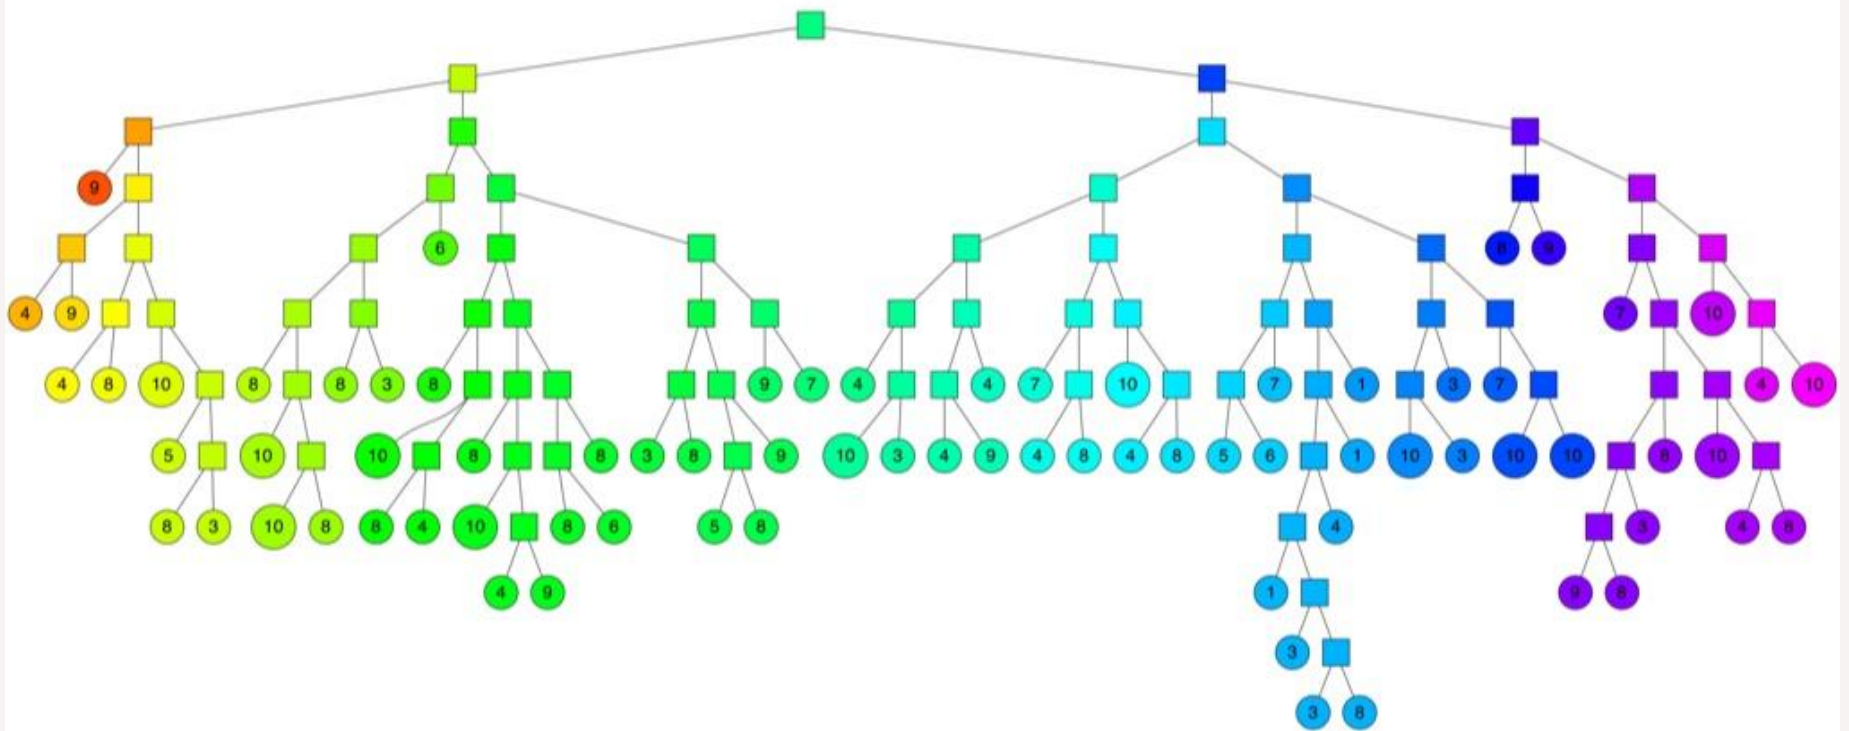


Again...





# Binary tree



```
from annoy import AnnoyIndex

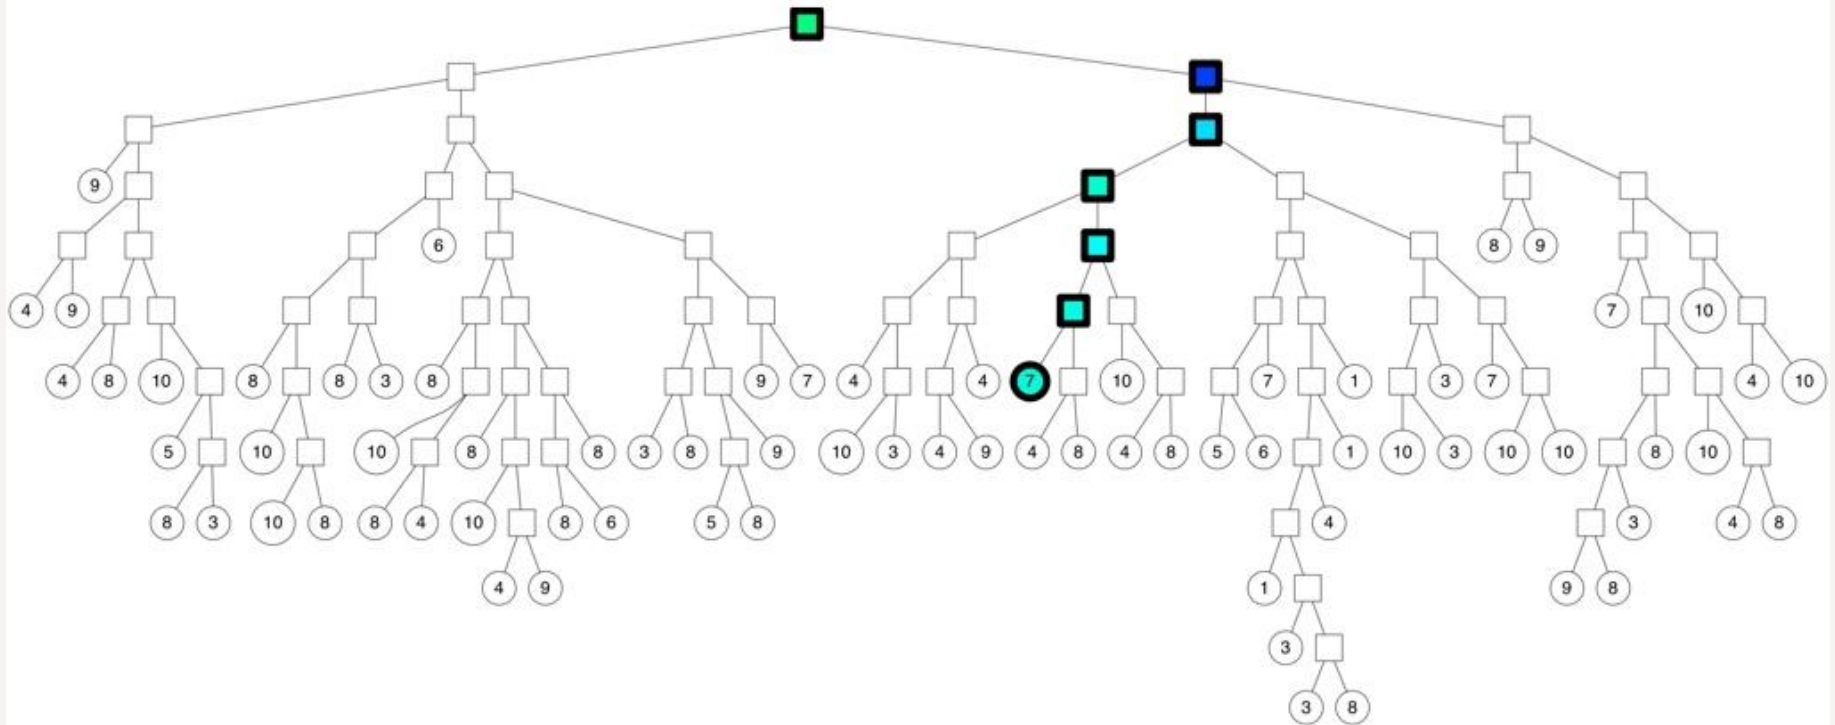
# init index with 40 dimension features
ann_index = AnnoyIndex(40)

# add data
for i, features in enumerate(db_features):
    ann_index.add_item(i, features)

# perform indexing with 10 trees
ann_index.build(10)

# save the resulting index into a bin file
ann_index.save('test.ann')
```

# Searching the tree





```
from annoy import AnnoyIndex

ann_index = AnnoyIndex(40)

# load saved index
ann_index.load('test.ann')

# return 10 nearest image indices
ann_index.get_nns_by_vector(query_features, 10)
```

Let's break down into 2 questions:

- 1. What kind of features to be used? ✓**
- 2. How to search similar features? ✓**

**OK Cool! Then what?**

**Clone project:**

<https://github.com/arifgodari/simple-image-search>

**Download Sample Dataset:**

<http://bitly.com/holiday-small>

Take home questions:

- 1. What features that important?**
- 2. How to search efficiently and effectively?**
- 3. When to use Euclidean distance vs Cosine Similarity?**
- 4. How to handle huge dimension?**

**“Searching is Easy, Finding is Not”**

A photograph of Tom Hanks from the chest up, wearing a dark blue suit, a light blue striped shirt, and a dark tie with a subtle pattern. He has his signature curly brown hair and is looking slightly to his right with a neutral expression. The background is a soft-focus outdoor scene with trees displaying autumn foliage in shades of orange and yellow. The entire image is framed by a solid light orange border.

**T. HANKS**